

▼ Linear supervised regression

0. Import library

Import library

```
1  # Import libraries
2
3  # math library
4  import numpy as np
5
6  # visualization library
7  %matplotlib inline
8  from IPython.display import set_matplotlib_formats
9  set_matplotlib_formats('png2x','pdf')
10 import matplotlib.pyplot as plt
11
12 # machine learning library
13 from sklearn.linear_model import LinearRegression
14
15 # 3d visualization
16 from mpl_toolkits.mplot3d import axes3d
17
18 # computational time
19 import time
20
```

▼ 1. Load dataset


Load a set of data pairs $\{x_i, y_i\}_{i=1}^n$ where x represents label and y represents target.

```
1  # import data with numpy
2  path = '/assignment_02_profit_population.txt'
3  data = np.loadtxt(path, delimiter=',')
```

▼ 2. Explore the dataset distribution

Plot the training data points.

```
1  x_train = data[:,0]
2  y_train = data[:,1]
3
4  fig_1 = plt.figure(1)
5  plt.title("Training data")
6  plt.xlabel("Population size (x 10$)")
7  plt.ylabel("Profit (x 10k$)")
8  plt.scatter(x_train, y_train, c="r", marker="x")
9  plt.show()
10 fig_1.savefig('training data.png')
```



3. Define the linear prediction function

$$f_w(x) = w_0 + w_1 x$$

Vectorized implementation:

$$f_w(x) = Xw$$

with

$$X = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix} \quad \text{and} \quad w = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} \quad \Rightarrow \quad f_w(x) = Xw = \begin{bmatrix} w_0 + w_1 x_1 \\ w_0 + w_1 x_2 \\ \vdots \\ w_0 + w_1 x_n \end{bmatrix}$$

Implement the vectorized version of the linear predictive function.

```

1  # construct data matrix
2  x0 = np.ones((x_train.reshape(-1,1).shape[0],1))
3  X = np.hstack((x0, x_train.reshape(-1,1)))
4
5  # parameters vector
6  w = np.ones(2)
7
8  # predictive function definition
9  def f_pred(X,w):
10
11      f = np.dot(X, w)
12
13      return f
14
15  # Test predictive function
16  y_pred = f_pred(X,w)

```

4. Define the linear regression loss

$$L(w) = \frac{1}{n} \sum_{i=1}^n \left(f_w(x_i) - y_i \right)^2$$

Vectorized implementation:

$$L(w) = \frac{1}{n}(Xw - y)^T(Xw - y)$$

with

$$Xw = \begin{bmatrix} w_0 + w_1x_1 \\ w_0 + w_1x_2 \\ \vdots \\ w_0 + w_1x_n \end{bmatrix} \quad \text{and} \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

Implement the vectorized version of the linear regression loss function.

```
1  # loss function definition
2  def loss_mse(y_pred,y):
3
4      loss = np.dot((y_pred - y).T, (y_pred - y)) / len(y)
5
6      return loss
7
8
9  # Test loss function
10 y = y_train # label
11 y_pred = f_pred(X,w) # prediction
12
13 loss = loss_mse(y_pred,y)
```

▼ 5. Define the gradient of the linear regression loss

Vectorized implementation: Given the loss

$$L(w) = \frac{1}{n}(Xw - y)^T(Xw - y)$$

The gradient is given by

$$\frac{\partial}{\partial w}L(w) = \frac{2}{n}X^T(Xw - y)$$

Implement the vectorized version of the gradient of the linear regression loss function.

```
1  # gradient function definition
2  def grad_loss(y_pred,y,X):
3      grad = 2 * np.dot(X.T, (y_pred - y)) / len(y)
4      return grad
5
6
7  # Test grad function
8  y_pred = f_pred(X,w)
9  grad = grad_loss(y_pred,y,X)
```

▼ 6. Implement the gradient descent algorithm

- Vectorized implementation:

$$w^{k+1} = w^k - \tau \frac{2}{n} X^T (Xw^k - y)$$

Implement the vectorized version of the gradient descent function.

Plot the loss values $L(w^k)$ with respect to iteration k the number of iterations.

```
1  # gradient descent function definition
2  def grad_desc(X, y, w_init, tau, max_iter):
3
4      L_iters = np.ones(max_iter) # record the loss values
5      w_iters = np.ones((max_iter, 2)) # record the parameter values
6      w = w_init # initialization
7
8      for i in range(max_iter): # loop over the iterations
9
10         y_pred = f_pred(X,w) # linear predication function
11         grad_f = grad_loss(y_pred,y,X) # gradient of the loss
12         w = w - (tau * grad_f) # update rule of gradient descent
13         L_iters[i] = loss_mse(y_pred,y) # save the current loss value
14         w_iters[i,:] = w # save the current w value
15
16     return w, L_iters, w_iters
17
18
19 # run gradient descent algorithm
20 start = time.time()
21 w_init = np.ones(2)
22 tau = 0.011
23 max_iter = 30
24
25 w, L_iters, w_iters = grad_desc(X,y,w_init,tau,max_iter)
26
27 print('Time=',time.time() - start) # plot the computational cost
28 print(L_iters[-1]) # plot the last value of the loss
29 print(w_iters[-1,:]) # plot the last value of the parameter w
30
31
32 # plot
33 fig_2 = plt.figure(2)
34 plt.plot(np.linspace(0, max_iter, max_iter), L_iters) # plot the loss curve
35 plt.xlabel('Iterations')
36 plt.ylabel('Loss')
37 plt.show()
38 fig_2.savefig('loss value.png')
```



Time= 0.0007340908050537109
 12.377896677532537
 [0.4235013 0.7596308]




▼ 7. Plot the linear prediction function

$$f_w(x) = w_0 + w_1 x$$

```

13 |
14 |
1 # linear regression model
2 x_pred = np.linspace(0,25,100) # define the domain of the prediction function
3 y_pred = w_iters[-1,0] + w_iters[-1,1] * x_pred # compute the prediction values within the given domain
4
5 # plot
6 fig_3 = plt.figure(3)
7 plt.scatter(x_train, y_train, c="r", marker="x")
8 plt.plot(x_pred, y_pred, c='b', label = "gradient descent optimization")
9 plt.legend(loc='best')
10 plt.title('Training data')
11 plt.xlabel('Population size (x 10k)')
12 plt.ylabel('Profit $(x 10k)')
13 plt.show()
14 fig_3.savefig('gradient descent optimization.png')

```



▼ 8. Comparison with Scikit-learn linear regression algorithm

Compare with the Scikit-learn solution

```

1 # run linear regression with scikit-learn
2 start = time.time()
3 lin_reg_sklearn = LinearRegression()
4 lin_reg_sklearn.fit(x_train.reshape(-1,1), y_train) # learn the model parameters
5 print('Time=',time.time() - start)
6
7
8 # compute loss value

```

```


9 w_sklearn = np.zeros([2,1])
10 w_sklearn[0,0] = lin_reg_sklearn.intercept_ # 절편
11 w_sklearn[1,0] = lin_reg_sklearn.coef_ # 기울기
12
13 print(w_sklearn)
14 y_pred_sklearn = lin_reg_sklearn.predict(x_train.reshape(-1,1)) # prediction obtained by the sklearn lib
15 loss_sklearn = loss_mse(y_pred_sklearn, y_train) # compute the loss from the sklearn solution
16
17 print('loss sklearn=',loss_sklearn)
18 print('loss gradient descent=',L_iters[-1])
19
20
21 # plot
22 fig_4 = plt.figure(4)
23 y_pred_skl = w_sklearn[0,0] + w_sklearn[1,0] * x_pred # compute the prediction values within the given d
24
25 plt.scatter(x_train.reshape(-1,1), y_train, c="r", marker="x")
26 plt.plot(x_pred, y_pred, c='b', label = "gradient descent optimization")
27 plt.plot(x_pred, y_pred_skl, c='r', label = "Scikit-learn optimization" )
28 plt.legend(loc='best')
29 plt.title('Training data')
30 plt.xlabel('Population size (x 10k)')
31 plt.ylabel('Profit $(x 10k)')
32 plt.show()
33 fig_4.savefig('scikit-learn optimization.png')

```

```

☞ Time= 0.025792837142944336
[[-3.89578088]
 [ 1.19303364]]
loss sklearn= 8.953942751950358
loss gradient descent= 12.377896677532537

```



9. Plot the loss surface, the contours of the loss and the gradient descent steps

```

1 # plot gradient descent
2 def plot_gradient_descent(X,y,w_init,tau,max_iter):
3     def f_pred(X,w):
4         f = np.dot(X, w)
5         return f
6

```

```

7     def loss_mse(y_pred,y):
8         loss = np.dot((y_pred - y).T, (y_pred - y)) / len(y)
9         return loss
10
11     # gradient descent function definition
12     def grad_desc(X, y, w_init, tau, max_iter):
13         L_iters = np.ones(max_iter) # record the loss values
14         w_iters = np.ones((max_iter, 2)) # record the parameter values
15         w = w_init # initialization
16
17         for i in range(max_iter): # loop over the iterations
18             y_pred = f_pred(X,w) # linear prediction function
19             grad_f = grad_loss(y_pred,y,X) # gradient of the loss
20             w = w - (tau * grad_f) # update rule of gradient descent
21             L_iters[i] = loss_mse(y_pred,y) # save the current loss value
22             w_iters[i,:] = w # save the current w value
23
24         return w, L_iters, w_iters
25
26     # run gradient descent
27     w, L_iters, w_iters = grad_desc(X, y, w_init, tau, max_iter)
28
29     # Create grid coordinates for plotting a range of L(w0,w1)-values
30     B0 = np.linspace(-10, 10, 50)
31     B1 = np.linspace(-1, 4, 50)
32
33     xx, yy = np.meshgrid(B0, B1, indexing='xy')
34     Z = np.zeros((B0.size,B1.size))
35
36     # Calculate loss values based on L(w0,w1)-values
37     for (i,j),v in np.ndenumerate(Z):
38         Z[i,j] = loss_mse(xx[i,j]+x_train*yy[i,j], y)
39
40     # 3D visualization
41     fig_5 = plt.figure(5)
42     fig_6 = plt.figure(6)
43     ax1 = fig_6.add_subplot()
44     ax2 = fig_5.add_subplot(projection='3d')
45
46     # Left plot
47     CS = ax1.contour(xx, yy, Z, np.logspace(-2, 3, 20), cmap=plt.cm.jet)
48     ax1.scatter(w[0], w[1], c="r")
49     ax1.plot(w_iters[:,0], w_iters[:,1])
50
51     # Right plot
52     ax2.plot_surface(xx, yy, Z, rstride=1, cstride=1, alpha=0.6, cmap=plt.cm.jet)
53     ax2.set_zlabel('Loss  $L(w_0,w_1)$ ')
54     ax2.set_zlim(Z.min(),Z.max())
55
56     # plot gradient descent
57     Z2 = np.zeros([max_iter])
58
59     for i in range(max_iter):
60         w0 = w_iters[i,0]
61         w1 = w_iters[i,1]
62         Z2[i] = L_iters[i]
63
64     ax2.plot(w_iters[:,0], w_iters[:,1], Z2)
65     ax2.scatter(w[0], w[1], c="r")

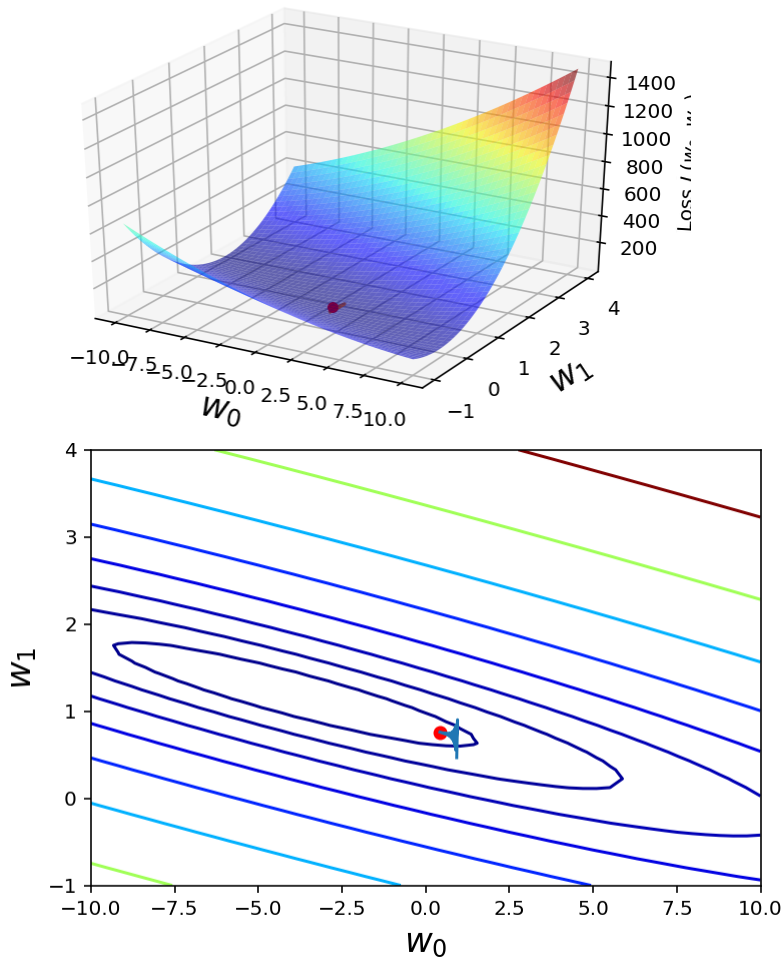
```

```

66
67     # settings common to both plots
68     ax1.set_xlabel(r'$w_0$', fontsize=17)
69     ax1.set_ylabel(r'$w_1$', fontsize=17)
70     ax2.set_xlabel(r'$w_0$', fontsize=17)
71     ax2.set_ylabel(r'$w_1$', fontsize=17)
72
73     return fig_5, fig_6

1  # run plot_gradient_descent function
2  w_init = np.ones(2)
3  tau = 0.011
4  max_iter = 30
5
6
7  fig_5, fig_6 = plot_gradient_descent(X,y,w_init,tau,max_iter)
8  fig_5.savefig('3D gradient.png')
9  fig_6.savefig('contour.png')


```



▼ Output results


▼ 1. Plot the training data (1pt)

1 fig_1




2. Plot the loss curve in the course of gradient descent (2pt)

1 fig_2




3. Plot the prediction function superimposed on the training data (2pt)

1 fig_3




4. Plot the prediction functions obtained by both the Scikit-learn linear regression solution and the gradient descent superimposed on the training data (2pt)

1 fig_4



5. Plot the loss surface (right) and the path of the gradient descent (2pt)

1 fig_5



6. Plot the contour of the loss surface (left) and the path of the gradient descent (2pt)

1 fig_6



