

▼ Supervised Logistic Regression for Classification

0. Import library

```
1  # Import libraries
2
3  # math library
4  import numpy as np
5
6  # visualization library
7  %matplotlib inline
8  from IPython.display import set_matplotlib_formats
9  set_matplotlib_formats('png2x','pdf')
10 import matplotlib.pyplot as plt
11
12 # machine learning library
13 from sklearn.linear_model import LogisticRegression
14
15 # 3d visualization
16 from mpl_toolkits.mplot3d import axes3d
17
18 # computational time
19 import time
20
```

▼ 1. Load dataset

The data features $x_i = (x_{i(1)}, x_{i(2)})$ represent 2 exam grades $x_{i(1)}$ and $x_{i(2)}$ for each student i .

The data label y_i indicates if the student i was admitted (value is 1) or rejected (value is 0).

```
1  # import data with numpy
2  data = np.loadtxt('/dataset.txt', delimiter=',')
3
4  # number of training data
5  n = data.shape[0]
6  print('Number of training data=',n)
```

➡ Number of training data= 100

▼ 2. Explore the dataset distribution

Plot the training data points.


You may use matplotlib function `scatter(x,y)`.

```
1  x1 = data[:,0] # exam grade 1
2  x2 = data[:,1] # exam grade 2
3  idx_admit = (data[:,2]==1) # index of students who were admitted
4  idx_rejec = (data[:,2]==0) # index of students who were rejected
5
6  fig_1 = plt.figure(figsize = (6,6))
7  plt.scatter(x1[idx_admit], x2[idx_admit], c= 'r', marker="+", label = 'Admitted')
8  plt.scatter(x1[idx_rejec], x2[idx_rejec], c= 'b', label = 'Rejected')
9  plt.title('Training data')
10 plt.xlabel('Exam grade 1')
11 plt.ylabel('Exam grade 2')
```

```

11 plt.ylabel('Exam grade 2')
12 plt.legend(loc = 'upper right')
13 plt.show()
14 fig_1.savefig('training data.png')

```



▼ 3. Sigmoid/logistic function

$$\sigma(\eta) = \frac{1}{1 + \exp^{-\eta}}$$

Define and plot the sigmoid function for values in [-10,10]:

You may use functions `np.exp`, `np.linspace`.

```

1 def sigmoid(z):
2
3     sigmoid_f = 1.0 / (1.0 + np.exp(-z))
4
5     return sigmoid_f
6
7
8 # plot
9 x_values = np.linspace(-10,10)
10
11 fig_2 = plt.figure(2)
12 plt.plot(x_values,sigmoid(x_values))
13 plt.title("Sigmoid function")
14 plt.grid(True)
15 fig_2.savefig('sigmoid function.png')

```



4. Define the prediction function for the classification

The prediction function is defined by:

$$p_w(x) = \sigma(w_0 + w_1x_{(1)} + w_2x_{(2)}) = \sigma(w^T x)$$

Implement the prediction function in a vectorised way as follows:

$$X = \begin{bmatrix} 1 & x_{1(1)} & x_{1(2)} \\ 1 & x_{2(1)} & x_{2(2)} \\ \vdots & & \\ 1 & x_{n(1)} & x_{n(2)} \end{bmatrix} \quad \text{and} \quad w = \begin{bmatrix} w_0 \\ w_1 \\ w_2 \end{bmatrix} \quad \Rightarrow \quad p_w(x) = \sigma(Xw) = \begin{bmatrix} \sigma(w_0 + w_1x_{1(1)} + w_2x_{1(2)}) \\ \sigma(w_0 + w_1x_{2(1)} + w_2x_{2(2)}) \\ \vdots \\ \sigma(w_0 + w_1x_{n(1)} + w_2x_{n(2)}) \end{bmatrix}$$

Use the new function `sigmoid`.

```

1  # construct the data matrix X
2  n = x1.shape[0]
3  X = np.ones([n,3])
4  X[:,1] = x1
5  X[:,2] = x2
6
7  # parameters vector
8  #w_init = np.array([-10, 0.1, -0.2])[:,None]
9
10 # predictive function definition
11 def f_pred(X,w):
12
13     p = sigmoid(np.dot(X,w))
14
15     return p
16
17 #y_pred = f_pred(X,w)
```

5. Define the classification loss function

Mean Square Error

$$L(w) = \frac{1}{n} \sum_{i=1}^n (\sigma(w^T x_i) - y_i)^2$$

Cross-Entropy

$$L(w) = \frac{1}{n} \sum_{i=1}^n (-y_i \log(\sigma(w^T x_i)) - (1 - y_i) \log(1 - \sigma(w^T x_i)))$$

The vectorized representation is for the mean square error is as follows:

$$L(w) = \frac{1}{n} (p_w(x) - y)^T (p_w(x) - y)$$

The vectorized representation is for the cross-entropy error is as follows:

$$L(w) = \frac{1}{n} (-y^T \log(p_w(x)) - (1 - y)^T \log(1 - p_w(x)))$$

where

$$p_w(x) = \sigma(Xw) = \begin{bmatrix} \sigma(w_0 + w_1x_{1(1)} + w_2x_{1(2)}) \\ \sigma(w_0 + w_1x_{2(1)} + w_2x_{2(2)}) \\ \vdots \\ \sigma(w_0 + w_1x_{n(1)} + w_2x_{n(2)}) \end{bmatrix} \quad \text{and} \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

You may use numpy functions `.T` and `np.log`.

```

1  def mse_loss(label, h_arr): # mean square error
2
3      return np.mean(np.square((h_arr - label)))
4
5  def ce_loss(label, h_arr): # cross-entropy error
6
7      return np.mean(-label * np.log(h_arr) - (1 - label) * np.log(1 - h_arr))

```

6. Define the gradient of the classification loss function

Given the mean square loss

$$L(w) = \frac{1}{n} (p_w(x) - y)^T (p_w(x) - y)$$

The gradient is given by

$$\frac{\partial}{\partial w} L(w) = \frac{2}{n} X^T ((p_w(x) - y) \odot (p_w(x) \odot (1 - p_w(x))))$$

Given the cross-entropy loss

$$L(w) = \frac{1}{n} (-y^T \log(p_w(x)) - (1 - y)^T \log(1 - p_w(x)))$$

The gradient is given by

$$\frac{\partial}{\partial w} L(w) = \frac{2}{n} X^T (p_w(x) - y)$$

Implement the vectorized version of the gradient of the classification loss function

```

1  # loss function definition
2  def loss_mse(y_pred,y):
3      n = len(y)
4      # diag(x)y.
5      loss = 2 * np.dot(X.T, (y_pred - y) * (y_pred * (1 - y_pred))) / n
6      return loss
7
8  def loss_ce(y_pred,y):
9      n = len(y)
10     loss = 2 * np.dot(X.T, (y_pred - y)) / n
11     return loss
12
13 # Test loss function
14 #y = data[:,2][:,None] # label
15 #y_pred = f_pred(X,w) # prediction
16
17 #loss = loss_mse(y_pred,y)

```

7. Implement the gradient descent algorithm

Vectorized implementation for the mean square loss:

$$w^{k+1} = w^k - \tau \frac{2}{n} X^T \left((p_w(x) - y) \odot (p_w(x) \odot (1 - p_w(x))) \right)$$

Vectorized implementation for the cross-entropy loss:

$$w^{k+1} = w^k - \tau \frac{2}{n} X^T (p_w(x) - y)$$

Plot the loss values $L(w^k)$ w.r.t. iteration k the number of iterations for the both loss functions.

```
1 # gradient descent function definition
2 def grad_desc_mse(X, y , w_init ,tau=1e-4, max_iter=500):
3     L_iters_mse = np.zeros([max_iter])
4     w_iters = np.zeros([max_iter,2]) # record the loss values
5     w = w_init # initialization
6
7     for i in range(max_iter): # loop over the iterations
8
9         y_pred = f_pred(X,w) # linear prediction function
10        grad_f = loss_mse(y_pred,y) # gradient of the loss
11        w = w - tau* grad_f # update rule of gradient descent
12        L_iters_mse[i] = mse_loss(y, y_pred) # save the current loss value
13
14        w_iters[i,:] = w[0],w[1] # save the current w value
15
16    return w, L_iters_mse, w_iters

```



```
1 # gradient descent function definition
2 def grad_desc_ce(X, y , w_init=np.array([1,1,1])[:,None] ,tau=1e-4, max_iter=500):
3     L_iters_ce = np.zeros([max_iter]) # record the loss values
4     w_iters = np.zeros([max_iter,2]) # record the loss values
5     w = w_init # initialization
6
7     for i in range(max_iter): # loop over the iterations
8
9         y_pred = f_pred(X,w) # linear prediction function
10        grad_f = loss_ce(y_pred,y) # gradient of the loss
11        w = w - tau* grad_f # update rule of gradient descent
12        L_iters_ce[i] = ce_loss(y, y_pred) # save the current loss value
13
14        w_iters[i,:] = w[0],w[1] # save the current w value
15
16    return w, L_iters_ce, w_iters

```

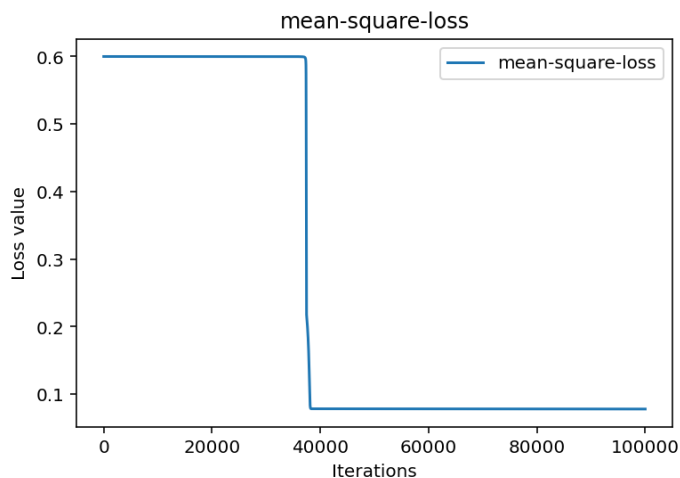


```
1 # run gradient descent algorithm
2 start = time.time()
3 w_init = np.array([-10, 0.1, -0.2])[:,None]
4 tau = 0.0005; max_iter = 100000
5
6 w_mse, L_iters_mse, w_iters_mse = grad_desc_mse(X,y,w_init,tau,max_iter)
7 # plot
8 fig_3 = plt.figure(3)
9 plt.plot(np.linspace(0, max_iter, max_iter), L_iters_mse, label = 'mean-square-loss')
10 print('mse loss:', L_iters_mse[-1])
11 print('mse weight:', w_mse)
12 plt.xlabel('Iterations')
13 plt.ylabel('Loss value')
14 plt.legend(loc = 'upper right')
15 plt.title('mean-square-loss')
16 plt.show()
17 fig_3.savefig('mse_loss.png')
```

```

↗ mse loss: 0.07804565596525387
mse weight: [[-10.10936997]
 [ 0.08670915]
 [ 0.07883981]]

```



```


1 # run gradient descent algorithm
2 start = time.time()
3 w_init = np.array([-10, 0.1, -0.2])[:,None]
4 tau = 0.0005; max_iter = 100000
5
6 w_ce, L_iters_ce, w_iters_ce = grad_desc_ce(X,y,w_init,tau,max_iter)
7 # plot
8 fig_4 = plt.figure(3)
9 plt.plot(np.linspace(0, max_iter, max_iter), L_iters_ce, label = 'cross-entropy-loss')
10 print('ce loss:', L_iters_ce[-1])
11 print('ce weight:', w_ce)
12 plt.xlabel('Iterations')
13 plt.ylabel('Loss value')
14 plt.legend(loc = 'upper right')
15 plt.title('cross-entropy-loss')
16 plt.show()
17 fig_4.savefig('ce_loss.png')

```

```

↗ ce loss: 0.2576879684547988
ce weight: [[-11.22782193]
 [ 0.09515914]
 [ 0.0890087 ]]

```




▼ 8. Plot the decision boundary

The decision boundary is defined by all points

$$x = (x_{(1)}, x_{(2)}) \quad \text{such that} \quad p_w(x) = 0.5$$

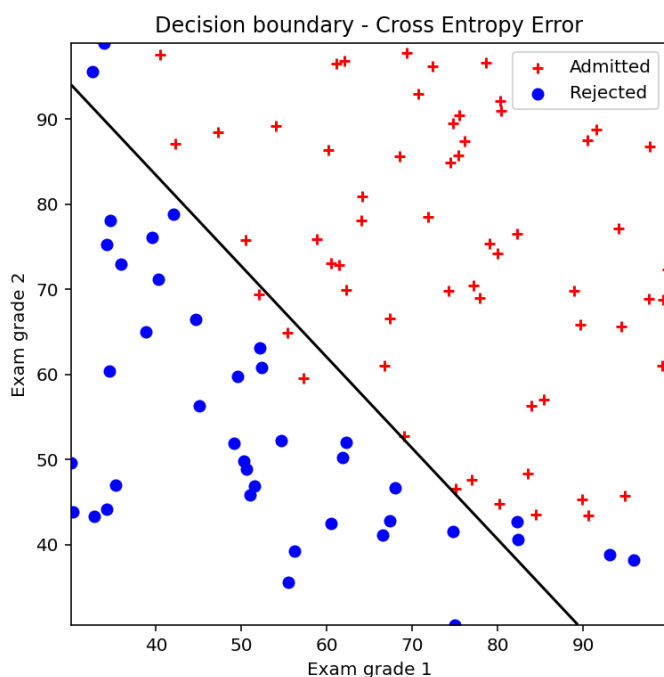
You may use numpy and matplotlib functions `np.meshgrid`, `np.linspace`, `reshape`, `contour`.

```
1 # compute values p(x) for multiple data points x
2 x1_min, x1_max = X[:,1].min(), X[:,1].max() # min and max of grade 1
3 x2_min, x2_max = X[:,2].min(), X[:,2].max() # min and max of grade 2
4 xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max), np.linspace(x2_min, x2_max)) # create meshgrid
5 X2 = np.ones([np.prod(xx1.shape),3])
6 X2[:,1] = xx1.reshape(-1)
7 X2[:,2] = xx2.reshape(-1)
8 p_mse = f_pred(X2,w_mse)
9 p_mse = p_mse.reshape(50,50)
10
11
12 # plot
13 fig_5 = plt.figure(5,figsize=(6,6))
14 plt.scatter(x1[idx_admit], x2[idx_admit], c= 'r', marker="+", label = 'Admitted')
15 plt.scatter(x1[idx_rejec], x2[idx_rejec], c= 'b', label = 'Rejected')
16 plt.contour(xx1, xx2, p_mse, [0.5], colors = 'k')
17 plt.xlabel('Exam grade 1')
18 plt.ylabel('Exam grade 2')
19 plt.legend(loc = 'upper right')
20 plt.title('Decision boundary - Mean square Error')
21 plt.show()
22 fig_5.savefig('MSE_boundary.png')
```



```
1 # compute values p(x) for multiple data points x
2 x1_min, x1_max = X[:,1].min(), X[:,1].max() # min and max of grade 1
3 x2_min, x2_max = X[:,2].min(), X[:,2].max() # min and max of grade 2
4 xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max), np.linspace(x2_min, x2_max)) # create meshgrid
5 X2 = np.ones([np.prod(xx1.shape),3])
6 X2[:,1] = xx1.reshape(-1)
7 X2[:,2] = xx2.reshape(-1)
8 p_ce = f_pred(X2,w_ce)
9 p_ce = p_ce.reshape(50,50)
10
11 # plot
12 fig_6 = plt.figure(4,figsize=(6,6))
13 plt.scatter(x1[idx_admit], x2[idx_admit], c= 'r', marker="+", label = 'Admitted')
14 plt.scatter(x1[idx_rejec], x2[idx_rejec], c= 'b', label = 'Rejected')
15 plt.contour(xx1, xx2, p_ce, [0.5], colors = 'k')
16 plt.xlabel('Exam grade 1')
17 plt.ylabel('Exam grade 2')
18 plt.legend(loc = 'upper right')
19 plt.title('Decision boundary - Cross Entropy Error')
20 plt.show()
```

```
plt.savefig('CE_boundary.png')
```



9. Comparison with Scikit-learn logistic regression algorithm with the gradient descent with the cross-entropy loss

You may use scikit-learn function `LogisticRegression(C=1e6)`.

```
1  # run logistic regression with scikit-learn
2  X0 = np.ones([n,2])
3  X0[:,0] = x1
4  X0[:,1] = x2
5
6  start = time.time()
7  logreg_sklern = LogisticRegression(solver = 'lbfgs', tol=0.0005, max_iter=10000) # scikit-learn logistic regression
8  logreg_sklern.fit(X0, y.reshape((y.shape[0],))) # learn the model parameters
9
10 # compute loss value
11 print(logreg_sklern.coef_)
12 w_sklern = np.array([1.,1.,1.])[1:,None]
13 w_sklern[0,0] = logreg_sklern.intercept_ # bias(w0)
14 w_sklern[1:3,0] = logreg_sklern.coef_[0,0:2].T # w1, w2
15
16 loss_sklern = loss_ce(logreg_sklern.predict(X0), y.reshape((y.shape[0],)))
17
18 # plot
19 fig_7 = plt.figure(4,figsize=(6,6))
20 plt.scatter(x1[idx_admit], x2[idx_admit], c = 'r', marker="+", label = 'Admitted')
21 plt.scatter(x1[idx_rejec], x2[idx_rejec], c = 'b', label = 'Rejected')
22 plt.xlabel('Exam grade 1')
23 plt.ylabel('Exam grade 2')
24
25 x1_min, x1_max = X[:,1].min(), X[:,1].max() # grade 1
26 x2_min, x2_max = X[:,2].min(), X[:,2].max() # grade 2
27
28 xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max), np.linspace(x2_min, x2_max)) # create meshgrid
29
30 X3 = np.ones([np.prod(xx1.shape),3])
31 X3[:,1] = xx1.reshape(-1)
32 X3[:,2] = xx2.reshape(-1)
33
34 p2 = f_pred(X3,w_sklern)
```




```

35 p2 = p2.reshape(50,50)
36 plt.contour(xx1, xx2, p2, [0.5], colors = 'r');
37
38 plt.title('Decision boundary (black with gradient descent and magenta with scikit-learn)')
39 plt.legend(loc = 'upper right')
40 plt.show()
41 fig_7.savefig('Decision boundary-sikit-learn.png')

```

[[0.20535491 0.2005838]]

Decision boundary (black with gradient descent and magenta with scikit-learn)



▼ 10. Plot the probability map

```

1 num_a = 110
2 grid_x1 = np.linspace(20,110,num_a)
3 grid_x2 = np.linspace(20,110,num_a)
4
5 score_x1, score_x2 = np.meshgrid(grid_x1, grid_x2)
6
7 Z = np.ones([np.prod(score_x1.shape) ,3])
8 Z[:,1] = score_x1.reshape(-1)
9 Z[:,2] = score_x2.reshape(-1)
10 predict_prob_mse = sigmoid(np.dot(Z,w_mse))
11 predict_prob_mse = predict_prob_mse.reshape(110,110)
12
13 predict_prob_ce = sigmoid(np.dot(Z,w_ce))
14 predict_prob_ce = predict_prob_ce.reshape(110,110)
15
16 #for i in range(len(score_x1)):
17     #for j in range(len(score_x2)):
18
19
20         #Z[j, i] = predict_prob
21
22         # actual plotting example

```

```


1 fig_8 = plt.figure(figsize=(10,10))
2
3 ax = fig_8.add_subplot(111)
4 ax.tick_params( )
5 ax.set_xlabel('Exam grade 1')
6 ax.set_ylabel('Exam grade 2')
7
8 ax.set_xlim(20, 110)

```

```

8 ax.set_xlim(20, 110)
9 ax.set_ylim(20, 110)
10
11 cf = ax.contourf(score_x1, score_x2, predict_prob_mse, 100, cmap = "bwr", vmin = 0, vmax = 1, alpha = 0.6)
12 ax.scatter(x1[idx_admit], x2[idx_admit], c= 'r', marker="+", label = 'Admitted')
13 ax.scatter(x1[idx_rejec], x2[idx_rejec], c= 'b', label = 'Rejected')
14 cbar = fig_8.colorbar(cf)
15 cbar.update_ticks()
16
17 plt.legend(loc = 'upper right')
18 plt.show()
19 fig_8.savefig('Probability Map.png')

```




```

1 fig_9 = plt.figure(figsize=(10,10))
2
3 ax = fig_9.add_subplot(111)
4 ax.tick_params( )
5 ax.set_xlabel('Exam grade 1')
6 ax.set_ylabel('Exam grade 2')
7
8 ax.set_xlim(20, 110)
9 ax.set_ylim(20, 110)
10
11 cf = ax.contourf(score_x1, score_x2, predict_prob_ce, 100, cmap = "bwr", vmin = 0, vmax = 1, alpha = 0.6)
12 ax.scatter(x1[idx_admit], x2[idx_admit], c= 'r', marker="+", label = 'Admitted')
13 ax.scatter(x1[idx_rejec], x2[idx_rejec], c= 'b', label = 'Rejected')
14 cbar = fig_9.colorbar(cf)
15 cbar.update_ticks()
16
17 plt.legend(loc = 'upper right')
18 plt.show()
19 fig_8.savefig('Probability Map.png')

```






▼ Output results


▼ 1. Plot the dataset in 2D cartesian coordinate system (1pt)

1 fig_1



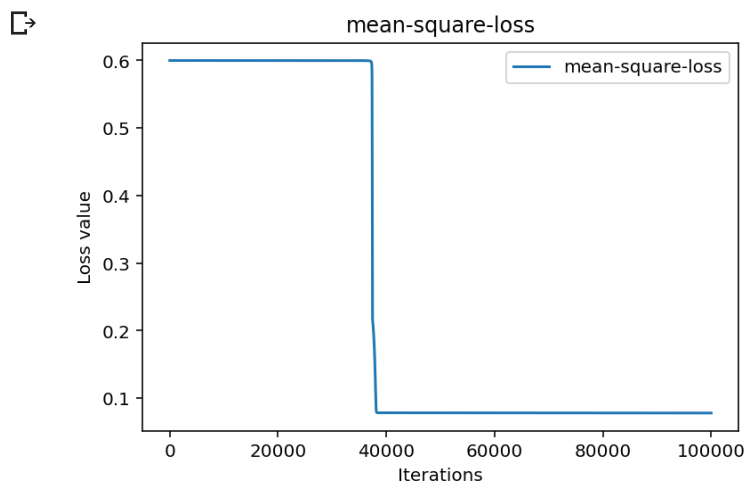
▼ 2. Plot the sigmoid function (1pt)

1 fig_2



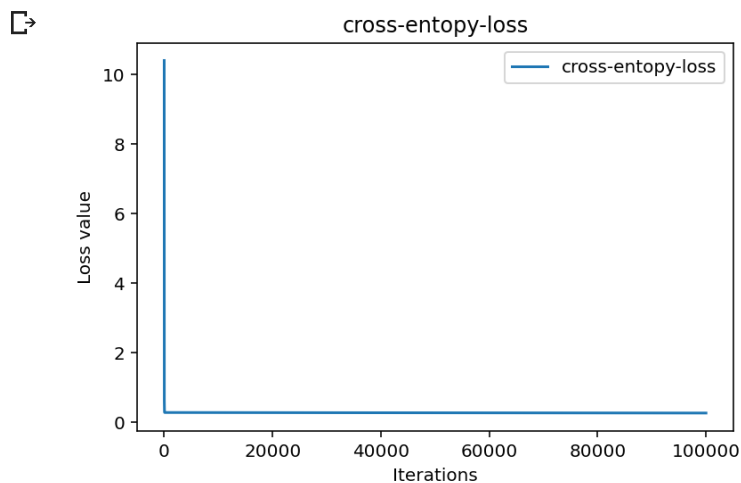
3. Plot the loss curve in the course of gradient descent using the mean square error (2pt)

1 fig_3




4. Plot the loss curve in the course of gradient descent using the cross-entropy error (2pt)

1 fig_4



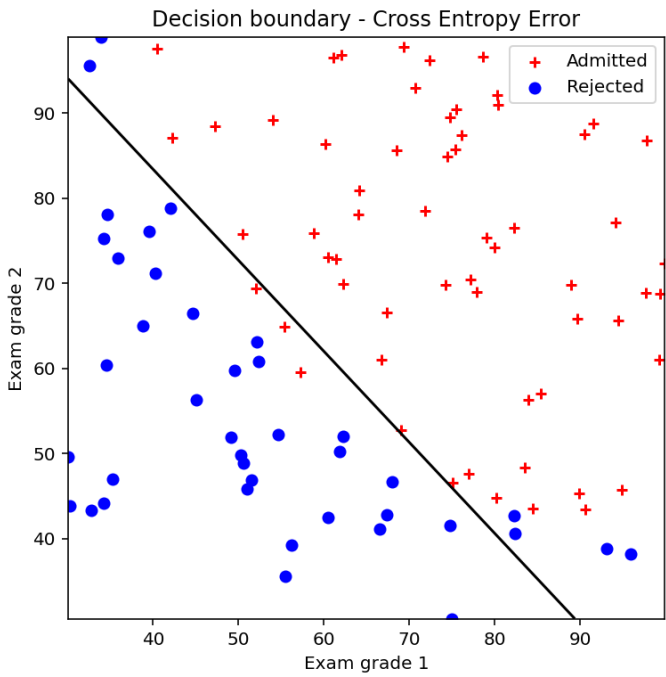
5. Plot the decision boundary using the mean square error (2pt)

1 fig_5



6. Plot the decision boundary using the cross-entropy error (2pt)

1 fig_6




7. Plot the decision boundary using the Scikit-learn logistic regression algorithm (2pt)

1 fig_7




Decision boundary (black with gradient descent and magenta with scikit-learn)



▼ 8. Plot the probability map using the mean square error (2pt)

1 fig_8



▼ 9. Plot the probability map using the cross-entropy error (2pt)

1 fig_9



