

# ▼ Convolutional Neural Network for the classification task on MNIST

Design a neural network that consists of a sequence of convolutional layers for characterizing features and a sequence of fully connected layers for classifying the characteristic features into categories.

```
1  import os
2  import math
3
4  # load data
5  from torch.utils.data import DataLoader
6  from torchvision import datasets, transforms
7
8  # train
9  import torch
10 from torch import nn, optim
11 from torch.nn import functional as F
12 from torch.optim import lr_scheduler
13 import numpy as np
14
15 # visualization
16 import matplotlib.pyplot as plt
17 import pandas as pd
```

check device

```
1  device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
2  print('Device: {}'.format(device))

Device: cuda:0
```

## ▼ 1. Data

- you can use any data normalisation method
- one example of the data normalisation is whitening as given by:

```
1  transform_train = transforms.Compose([
2      transforms.ToTensor(),
3      transforms.Normalize(mean=(0.5,), std=(0.5,))
4  ])
5
6  transform_test = transforms.Compose([
7      transforms.ToTensor(),
8      transforms.Normalize(mean=(0.5,), std=(0.5,))
9  ])
10
```

- load the MNIST dataset
- use the original training dataset for testing your model
- use the original testing dataset for training your model

```
1  data_path = './MNIST'
2
```

```

2
3 data_test = datasets.MNIST(root = data_path, train= True, download=True, transform=
4 data_train = datasets.MNIST(root = data_path, train= False, download=True, transform=

```

- Note that the number of your training data must be 10,000
- Note that the number of your testing data must be 60,000

```

1 print("the number of your training data (must be 10,000) = ", data_train.__len__())
2 print("hte number of your testing data (must be 60,000) = ", data_test.__len__())

```

```

the number of your training data (must be 10,000) = 10000
hte number of your testing data (must be 60,000) = 60000

```

## 2. Model

- design a neural network architecture with a combination of convolutional layers and fully connected layers
- use any number of feature layers (convolutional layers)
- use any size of convolutional kernel\_size
- use any dimension of classification layers
- use any type of activation functions
- one example model of the convolutional neural network is as follows:

```

1 class MyModel(nn.Module):
2     def __init__(self):
3         super(MyModel, self).__init__()
4
5         # feature layer
6         self.features = nn.Sequential(
7             nn.Conv2d(1, 32, kernel_size=3, stride=1, padding=1),
8             nn.BatchNorm2d(32),
9             nn.ReLU(inplace=True),
10
11             nn.Conv2d(32, 32, kernel_size=3, stride=1, padding=1),
12             nn.BatchNorm2d(32),
13             nn.ReLU(inplace=True),
14             nn.MaxPool2d(kernel_size=2, stride=2),
15
16             nn.Conv2d(32, 64, kernel_size=3, padding=1),
17             nn.BatchNorm2d(64),
18             nn.ReLU(inplace=True),
19
20             nn.Conv2d(64, 64, kernel_size=3, padding=1),
21             nn.BatchNorm2d(64),
22             nn.ReLU(inplace=True),
23             nn.MaxPool2d(kernel_size=2, stride=2)
24         )
25
26         # classifier layer
27         self.classifier = nn.Sequential(
28             nn.Dropout(p = 0.5),
29             nn.Linear(64 * 7 * 7, 512),
30             nn.BatchNorm1d(512),
31             nn.ReLU(inplace=True),
32
33             nn.Dropout(p = 0.5),
34             nn.Linear(512, 512),
35             nn.BatchNorm1d(512)

```

```

36         nn.BatchNorm1d(512),
37         nn.ReLU(inplace=True),
38         nn.Dropout(p = 0.5),
39         nn.Linear(512, 10),
40     )
41
42
43     for m in self.features.children():
44         if isinstance(m, nn.Conv2d):
45             n = m.kernel_size[0] * m.kernel_size[1] * m.out_channels
46             m.weight.data.normal_(0, math.sqrt(2. / n))
47         elif isinstance(m, nn.BatchNorm2d):
48             m.weight.data.fill_(1)
49             m.bias.data.zero_()
50
51     for m in self.classifier.children():
52         if isinstance(m, nn.Linear):
53             nn.init.xavier_uniform_(m.weight)
54         elif isinstance(m, nn.BatchNorm1d):
55             m.weight.data.fill_(1)
56             m.bias.data.zero_()
57
58     def forward(self, x):
59         x = self.features(x)
60         x = x.view(x.size(0), -1)
61         x = self.classifier(x)
62
63     return x

```

### 3. Loss function

- use any type of loss function
- design the output of the output layer considering your loss function

```
1 criterion = nn.CrossEntropyLoss()
```

### 4. Optimization

- use any stochastic gradient descent algorithm for the optimization
- use any size of the mini-batch
- use any optimization algorithm (for example, Momentum, AdaGrad, RMSProp, Adam)
- use any regularization algorithm (for example, Dropout, Weight Decay)
- use any annealing scheme for the learning rate (for example, constant, decay, staircase)

```
1 BATCH_SIZE = 32
```

```
1 train_loader = torch.utils.data.DataLoader(data_train, batch_size=BATCH_SIZE, shuffle=True)
2 test_loader = torch.utils.data.DataLoader(data_test, batch_size=BATCH_SIZE, shuffle=False)
```

```
1 model = MyModel()
2 model.to(device)
```

```

MyModel(
  (features): Sequential(
    (0): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (4): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
    (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (7): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (9): ReLU(inplace=True)
    (10): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (12): ReLU(inplace=True)
    (13): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=3136, out_features=512, bias=True)
    (2): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (3): ReLU(inplace=True)
    (4): Dropout(p=0.5, inplace=False)
    (5): Linear(in_features=512, out_features=512, bias=True)
    (6): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (7): ReLU(inplace=True)
    (8): Dropout(p=0.5, inplace=False)
    (9): Linear(in_features=512, out_features=10, bias=True)
  )
)

```

```

1  epochs = 100
2  lr = 0.003
3  step_size = 7
4
5  optimizer = optim.Adam(model.parameters(), lr=lr)
6  exp_lr_scheduler = lr_scheduler.StepLR(optimizer, step_size=step_size, gamma=0.1)

```

## 5. Training

```

1  test_loss_min = np.Inf
2  train_losses = []
3  test_losses = []
4  history_accuracy = []
5  history_running_acc = []
6
7  for e in range(1, epochs+1):
8      running_loss = 0
9      running_acc = 0
10
11     for images, labels in train_loader:
12         model.train()
13         images, labels = images.to(device), labels.to(device)
14
15         optimizer.zero_grad()
16         ps = model(images)
17         _, top_class = ps.topk(1, dim=1)
18         equals = top_class == labels.view(*top_class.shape)
19
20         loss = criterion(ps, labels)
21         running_loss += torch.mean(equals.type(torch.FloatTensor))

```

```
running_acc += torch.mean(equals.type(torch.FloatTensor))
```

```
loss.backward()  
optimizer.step()
```

```
running_loss += loss.item()
```

```
else:
```

```
test_loss = 0  
accuracy = 0
```

```
with torch.no_grad():
```

```
    model.eval()  
    for images, labels in test_loader:  
        images, labels = images.to(device), labels.to(device)
```

```
        ps = model(images)  
        _, top_class = ps.topk(1, dim=1)  
        equals = top_class == labels.view(*top_class.shape)
```

```
        test_loss += criterion(ps, labels).item()  
        accuracy += torch.mean(equals.type(torch.FloatTensor))
```

```
train_losses.append(running_loss/len(train_loader))  
test_losses.append(test_loss/len(test_loader))  
history_accuracy.append(accuracy/len(test_loader))  
history_running_acc.append(running_acc/len(train_loader))
```

```
if (accuracy/len(test_loader) > 0.9893):  
    break  
exp_lr_scheduler.step()
```

```
print(f"Epoch: {e}/{epochs}.. ",  
      f"Training Loss: {running_loss/len(train_loader):.3f}.. ",  
      f"Testing Loss: {test_loss/len(test_loader):.3f} / ",  
      f"Train Accuracy: {running_acc/len(train_loader):.3f} ",  
      f"Test Accuracy: {accuracy/len(test_loader):.3f}")
```

Epoch: 1/100..	Training Loss: 0.393..	Testing Loss: 0.106	/	Train Accuracy: 0.879	Test Accuracy: 0.967
Epoch: 2/100..	Training Loss: 0.168..	Testing Loss: 0.069	/	Train Accuracy: 0.948	Test Accuracy: 0.980
Epoch: 3/100..	Training Loss: 0.124..	Testing Loss: 0.067	/	Train Accuracy: 0.961	Test Accuracy: 0.979
Epoch: 4/100..	Training Loss: 0.109..	Testing Loss: 0.068	/	Train Accuracy: 0.966	Test Accuracy: 0.980
Epoch: 5/100..	Training Loss: 0.100..	Testing Loss: 0.076	/	Train Accuracy: 0.971	Test Accuracy: 0.977
Epoch: 6/100..	Training Loss: 0.084..	Testing Loss: 0.049	/	Train Accuracy: 0.975	Test Accuracy: 0.985
Epoch: 7/100..	Training Loss: 0.068..	Testing Loss: 0.058	/	Train Accuracy: 0.978	Test Accuracy: 0.983
Epoch: 8/100..	Training Loss: 0.061..	Testing Loss: 0.047	/	Train Accuracy: 0.982	Test Accuracy: 0.986
Epoch: 9/100..	Training Loss: 0.044..	Testing Loss: 0.046	/	Train Accuracy: 0.986	Test Accuracy: 0.987
Epoch: 10/100..	Training Loss: 0.039..	Testing Loss: 0.042	/	Train Accuracy: 0.987	Test Accuracy: 0.988
Epoch: 11/100..	Training Loss: 0.039..	Testing Loss: 0.040	/	Train Accuracy: 0.986	Test Accuracy: 0.988
Epoch: 12/100..	Training Loss: 0.037..	Testing Loss: 0.043	/	Train Accuracy: 0.988	Test Accuracy: 0.988
Epoch: 13/100..	Training Loss: 0.034..	Testing Loss: 0.040	/	Train Accuracy: 0.988	Test Accuracy: 0.988
Epoch: 14/100..	Training Loss: 0.026..	Testing Loss: 0.041	/	Train Accuracy: 0.991	Test Accuracy: 0.988
Epoch: 15/100..	Training Loss: 0.027..	Testing Loss: 0.040	/	Train Accuracy: 0.991	Test Accuracy: 0.988
Epoch: 16/100..	Training Loss: 0.025..	Testing Loss: 0.039	/	Train Accuracy: 0.993	Test Accuracy: 0.988
Epoch: 17/100..	Training Loss: 0.025..	Testing Loss: 0.039	/	Train Accuracy: 0.992	Test Accuracy: 0.988
Epoch: 18/100..	Training Loss: 0.029..	Testing Loss: 0.040	/	Train Accuracy: 0.991	Test Accuracy: 0.988
Epoch: 19/100..	Training Loss: 0.030..	Testing Loss: 0.040	/	Train Accuracy: 0.990	Test Accuracy: 0.988
Epoch: 20/100..	Training Loss: 0.025..	Testing Loss: 0.039	/	Train Accuracy: 0.992	Test Accuracy: 0.988
Epoch: 21/100..	Training Loss: 0.022..	Testing Loss: 0.038	/	Train Accuracy: 0.993	Test Accuracy: 0.988
Epoch: 22/100..	Training Loss: 0.026..	Testing Loss: 0.040	/	Train Accuracy: 0.991	Test Accuracy: 0.988
Epoch: 23/100..	Training Loss: 0.023..	Testing Loss: 0.038	/	Train Accuracy: 0.993	Test Accuracy: 0.988
Epoch: 24/100..	Training Loss: 0.022..	Testing Loss: 0.038	/	Train Accuracy: 0.993	Test Accuracy: 0.988
Epoch: 25/100..	Training Loss: 0.024..	Testing Loss: 0.039	/	Train Accuracy: 0.992	Test Accuracy: 0.988

Epoch: 26/100..	Training Loss: 0.023..	Testing Loss: 0.040	/	Train Accuracy: 0.993	Test Accuracy: 0.989
Epoch: 27/100..	Training Loss: 0.024..	Testing Loss: 0.038	/	Train Accuracy: 0.992	Test Accuracy: 0.989
Epoch: 28/100..	Training Loss: 0.023..	Testing Loss: 0.039	/	Train Accuracy: 0.992	Test Accuracy: 0.989
Epoch: 29/100..	Training Loss: 0.021..	Testing Loss: 0.039	/	Train Accuracy: 0.994	Test Accuracy: 0.989
Epoch: 30/100..	Training Loss: 0.026..	Testing Loss: 0.039	/	Train Accuracy: 0.991	Test Accuracy: 0.989
Epoch: 31/100..	Training Loss: 0.022..	Testing Loss: 0.040	/	Train Accuracy: 0.992	Test Accuracy: 0.989
Epoch: 32/100..	Training Loss: 0.023..	Testing Loss: 0.039	/	Train Accuracy: 0.992	Test Accuracy: 0.989




```
1 idx = (np.array(history_accuracy)[:>= 0.989])
2 print(np.array(history_accuracy)[idx])

[0.98906666 0.98908335 0.98915      0.9892667  0.98903334 0.98936665]
```

## 6. Visualization

### 1. Plot the training and testing losses over epochs [2pt]

```
1 fig_1 = plt.figure(figsize=(8,8))
2 plt.plot(np.array(range(len(train_losses))), train_losses, c = 'r', label = 'Training L
3 plt.plot(np.array(range(len(test_losses))), test_losses, c = 'b', label = 'Test Loss')
4 plt.legend(loc = 'upper right')
5 plt.title('Plot the loss curve')
6 plt.show()
7 fig_1.savefig('loss curve.png')
```



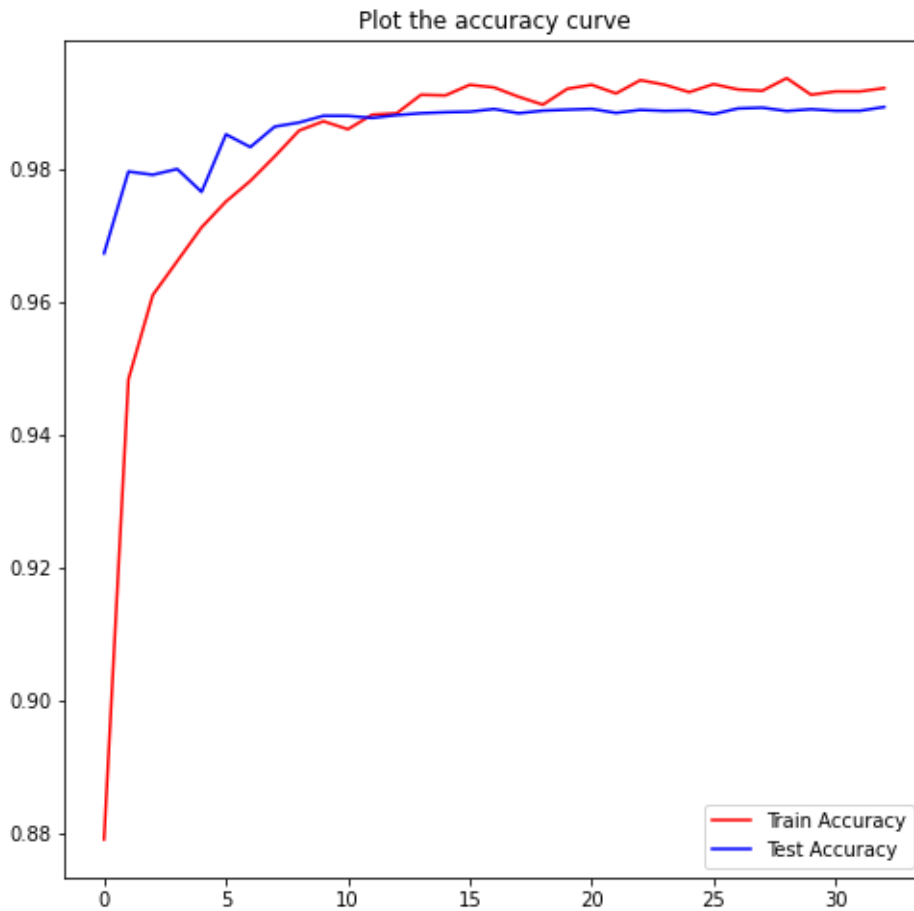
### 2. Plot the training and testing accuracies over epochs [2pt]

```
1 fig_2 = plt.figure(figsize=(8,8))
```

```

2 plt.plot(np.array(range(len(history_running_acc))), history_running_acc, c = 'r', label = 'Train Accuracy')
3 plt.plot(np.array(range(len(history_accuracy))), history_accuracy, c = 'b', label = 'Test Accuracy')
4 plt.legend(loc = 'lower right')
5 plt.title('Plot the accuracy curve')
6 plt.show()
7 fig_2.savefig('accuracy curve.png')

```



3. Print the final training and testing losses at convergence [2pt]

**\* NOTE \*** The values should be presented up to 5 decimal places

```

1 pd.options.display.float_format = '{:.5f}'.format
2 result_loss = pd.DataFrame({'loss': [train_losses[-1], test_losses[-1]]}, index = ['train', 'test'])
3 result_loss

```

	loss
training loss	0.02488
testing loss	0.03749

4. Print the final training and testing accuracies at convergence [20pt]

**\* NOTE \*** The values should be presented up to 5 decimal places (소수점 5째 자리까지 표기하시오)

```

1 result_acc = pd.DataFrame({'accuracy': [history_running_acc[-1].item(), history_accuracy[-1].item()]}, index = ['train', 'test'])
2 result_acc

```

## 5. Print the testing accuracies within the last 10 epochs [5pt]

**\* NOTE \*** The values should be presented up to 5 decimal places

```
1 print('>> Print the testing accuracies within the last 10 epochs [5pt]')
2 for e in range(10):
3     print(f"[epoch = {len(history_accuracy) - e}] {history_accuracy[-(e+1)]:.5f}")
```

```
>> Print the testing accuracies within the last 10 epochs [5pt]
```

```
[epoch = 33] 0.98937
```

```
[epoch = 32] 0.98880
```

```
[epoch = 31] 0.98880
```

```
[epoch = 30] 0.98903
```

```
[epoch = 29] 0.98877
```

```
[epoch = 28] 0.98927
```

```
[epoch = 27] 0.98915
```

```
[epoch = 26] 0.98830
```


```
[epoch = 25] 0.98885
```

```
[epoch = 24] 0.98878
```

## Submission

### 1. Plot the training and testing losses over epochs [2pt]


```
1 fig_1
```



### 2. Plot the training and testing accuracies over epochs [2pt]



1 fig\_2



3. Print the final training and testing losses at convergence [2pt]

1 result\_loss

	loss
<b>training loss</b>	0.02488
<b>testing loss</b>	0.03749

4. Print the final training and testing accuracies at convergence [20pt]

1 result\_acc

	accuracy
<b>training accuracy</b>	0.99221
<b>testing accuracy</b>	0.98937

5. Print the testing accuracies within the last 10 epochs [5pt]

```
1 print('>> Print the testing accuracies within the last 10 epochs [5pt]')
2 for e in range(10):
3     print(f"epoch = {len(history_accuracy) - e} {history_accuracy[-(e+1)]:.5f}")
```

```
>> Print the testing accuracies within the last 10 epochs [5pt]
[epoch = 33] 0.98937
[epoch = 32] 0.98880
[epoch = 31] 0.98880
[epoch = 30] 0.98903
[epoch = 29] 0.98877
[epoch = 28] 0.98927
[epoch = 27] 0.98915
[epoch = 26] 0.98830
[epoch = 25] 0.98885
[epoch = 24] 0.98878
```